

Faisal Baig

US CITIZEN

☎ 678-770-5117 | ✉ me@faisalbaig.com | 🏠 faisab.github.io | 🔗 linkedin.com/in/baig-faisal

Education

Georgia Institute of Technology

M.S. IN COMPUTER SCIENCE

- Graduate Student Assistant for CSE 6104 (Graduate CS Algorithms) - Wrote lecture notes & facilitated exams

Atlanta, Georgia

Grad December 2024

Georgia Institute of Technology

B.S. IN COMPUTER SCIENCE – ZELL MILLER SCHOLAR + DEANS LIST

Atlanta, Georgia

Grad May 2023

Experience

Boston Consulting Group

AI ENGINEER INTERN

Los Angeles, California

Jun. 2023 – Aug. 2023

- Incorporated multiple APIs within ingestion pipeline & built unified functionality for processing multiple formats: earnings calls, 10Ks, 10Qs
- Developed custom parsing functions for multiple formats, creating contextually richer content within token constraints (Python, LangChain, NLTK)
- Wrote ingestion functions for raw & semi-processed input data to ensure compatibility with tokenizer & vector DB
- Implemented data quality checks, filtering criteria, & post-processing for LLM output enhancing quality and reliability in user experience
- Created unit & integration tests to validate functionality and performance of data processing components (Python, Pytest, Docker, Terraform)

CertiK

PLATFORM ENGINEER INTERN

New York, New York

Sep. 2022 – Dec. 2022

- Automated 10+ anomaly detection models using Databricks/PySpark/Snowflake, implementing access controls, database integration, and model performance metrics logging with Datadog and MLflow
- Developed automated pipeline for creating & updating monitors of current data platform using Github Actions/Terraform. Automated migration of 320+ existing agents/monitors to infra-as-code (MonteCarlo/Datadog/Github Actions/Terraform)
- Deployed Terraform-based infrastructure to dev, staging, & live environments using Github Actions CI/CD pipelines for public and internal APIs and set up logging and monitoring (ECS, S3, Docker, Airflow, Terraform)
- Policy creation, environment setup, & IAM for ML Infrastructure and data pipelines (Databricks, Github, Snowflake)
- Investigated economic impact and proposed optimization for high consumption or high runtime Snowpipe loads & Snowflake queries

Equifax

SOFTWARE ENGINEER INTERN

Atlanta, Georgia

May 2022 – Aug. 2022

- Endpoint development and UI design of B2B portal using Springboot and Canva for the MVP and beyond of a new project (Java/BigQuery)
- Wrote functions to integrate novel data format in an active Beam ETL pipeline processing millions of records(Dataflow/Apache Beam/PostgreSQL)
- Conducted topic modeling with LDA on over 4 million support cases, top results populating unknown causes field (Python/Gensim/NLTK)
- Performed data exploration on support case dataset, reducing the feature set from 100+ columns to 13 key attributes (Python/Pandas)

Research

Georgia Tech AI Research and Discovery Lab

UNDERGRAD RESEARCHER

Atlanta, Georgia

Jan. 2021 - May 2022

- Developed full-stack app using Azure Blob Storage & Node.js serving REST API to allow categorization of CT Scans by radiologists at Emory University
- Test & log metrics for ML models trained on novel data to allow for categorization, localization, and gradation of different conditions lung conditions

Projects

IMAGE PROCESSING & QUALITY ASSESSMENT PIPELINE | PYTHON, OPENCV, SKIMAGE

- Multi-step image processing for consistent quality metrics: sharpness calculation, noise estimation, distortion measurement, and color accuracy.

ENSEMBLE LEARNING TRADING ALGORITHM | PYTORCH, SCIKIT-LEARN, REST

- Ensembles multiple supervised methods to maximize profit on historical trading data. Relevant features & stock groupings using K-Means w/DTW

PARALLELIZED MACHINE LEARNING LIBRARY | C++, OPENMP, EIGEN

- Parallelized algorithms (Linear Regression & K-means Clustering) and preprocessing in C++, leveraging Eigen for optimized matrix operations

Skills

Languages & Frameworks: Python, C++, JavaScript, SQL, Golang, PyTorch, PySpark, scikit-learn, OpenMP, d3.js

Tools: Git, AWS, Kubernetes, Docker, CI/CD, DataDog, Databricks, LangChain, NLTK, Microservices, Apache Airflow, Apache Beam, Jupyter, MLflow, Terraform, Make, Bash, Linux

Coursework: DS&A, Algorithm Analysis, Computer Architecture, Operating Systems, Software Design, Databases, Statistics, (MS) Machine Learning, (MS) Deep Learning, (MS) Network Security, (MS) High Performance Computing, (MS) Natural Language, (MS) Visual Analytics